

gem5-Garnet

Syam Sankar

Ph.D Scholar (Prime Minister's Research Fellow)
MARS Lab, Dept. of CSE, IIT Guwahati

- The default network model in Ruby is the **simple** network.
 - hop-by-hop network traversal (switches and links)
 - monitoring the available buffers and available bandwidth in output links before sending

Related Files:

- `src/mem/ruby/network/Network.py`
- `src/mem/ruby/network/simple`
- `src/mem/ruby/network/simple/SimpleNetwork.py`

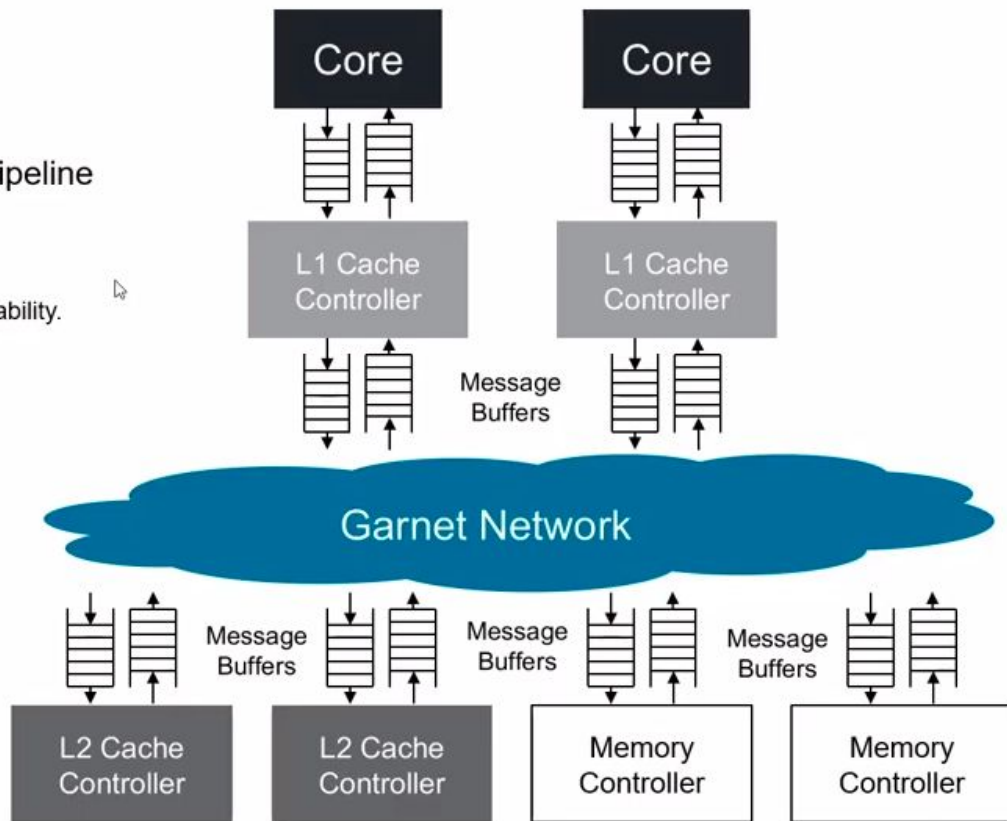
Garnet Network Simulator

Detailed NoC Model

- Part of Ruby Memory System in gem5
- Garnet 1.0 released in 2009: Fixed 5-stage router pipeline
- Garnet 2.0 released in 2016: Configurable routers
- Garnet 3.0 includes HeteroGarnet
 - Supports diverse heterogeneous networks with complete configurability.

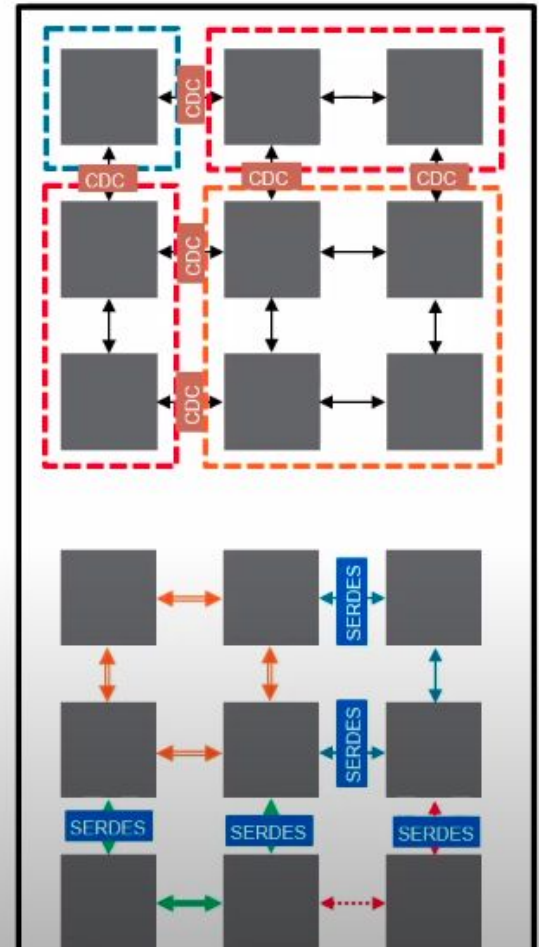
Features of Garnet 2.0

- Configurable topology
- Detailed router pipeline model
- Credit-based flow control
- Configurable routing policies
- Synthetic traffic simulations



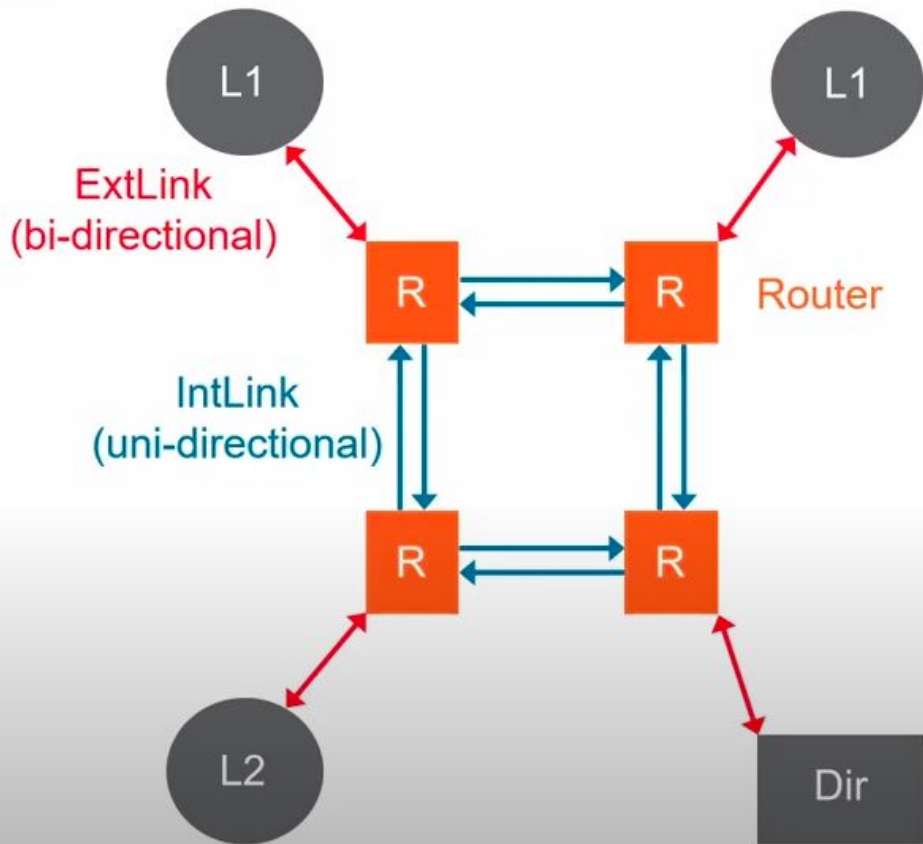
HeteroGarnet – Enabling Heterogeneous Networks

- Multiple clock domains
 - Each building block within network can be configured to be in a clock domain.
- Heterogeneous link widths
 - Each routers and link can be configured to support a distinct flit width
- Multiple sub-networks
 - Multiple mutually-exclusive networks could co-exist in a simulation
- Clock-domain crossing units
- Serializer-deserializer units



Building Topology in Garnet/HeteroGarnet


- Each topology is defined using a python file located in `configs/topologies`
- Building blocks on Garnet
 - Router**: Detailed router model with crossbar
 - ExtLink**: Used to connect to external controllers
 - IntLink**: Used to connect routers within network
- Each block has configurable parameters.



HeteroGarnet - CDC and SerDes Units


Clock Domain Crossing (CDC)

- Enables transmission of flits across clock domains
- Supports flow-control credit crossings
- Configurable latency
 - Static or Dynamic (Function of clock domain frequencies)
- Enabled by setting flags on the links
 - Internal Links: `src_cdc` and/or `dst_cdc`
 - External Links: `ext_cdc` and/or `int_cdc`

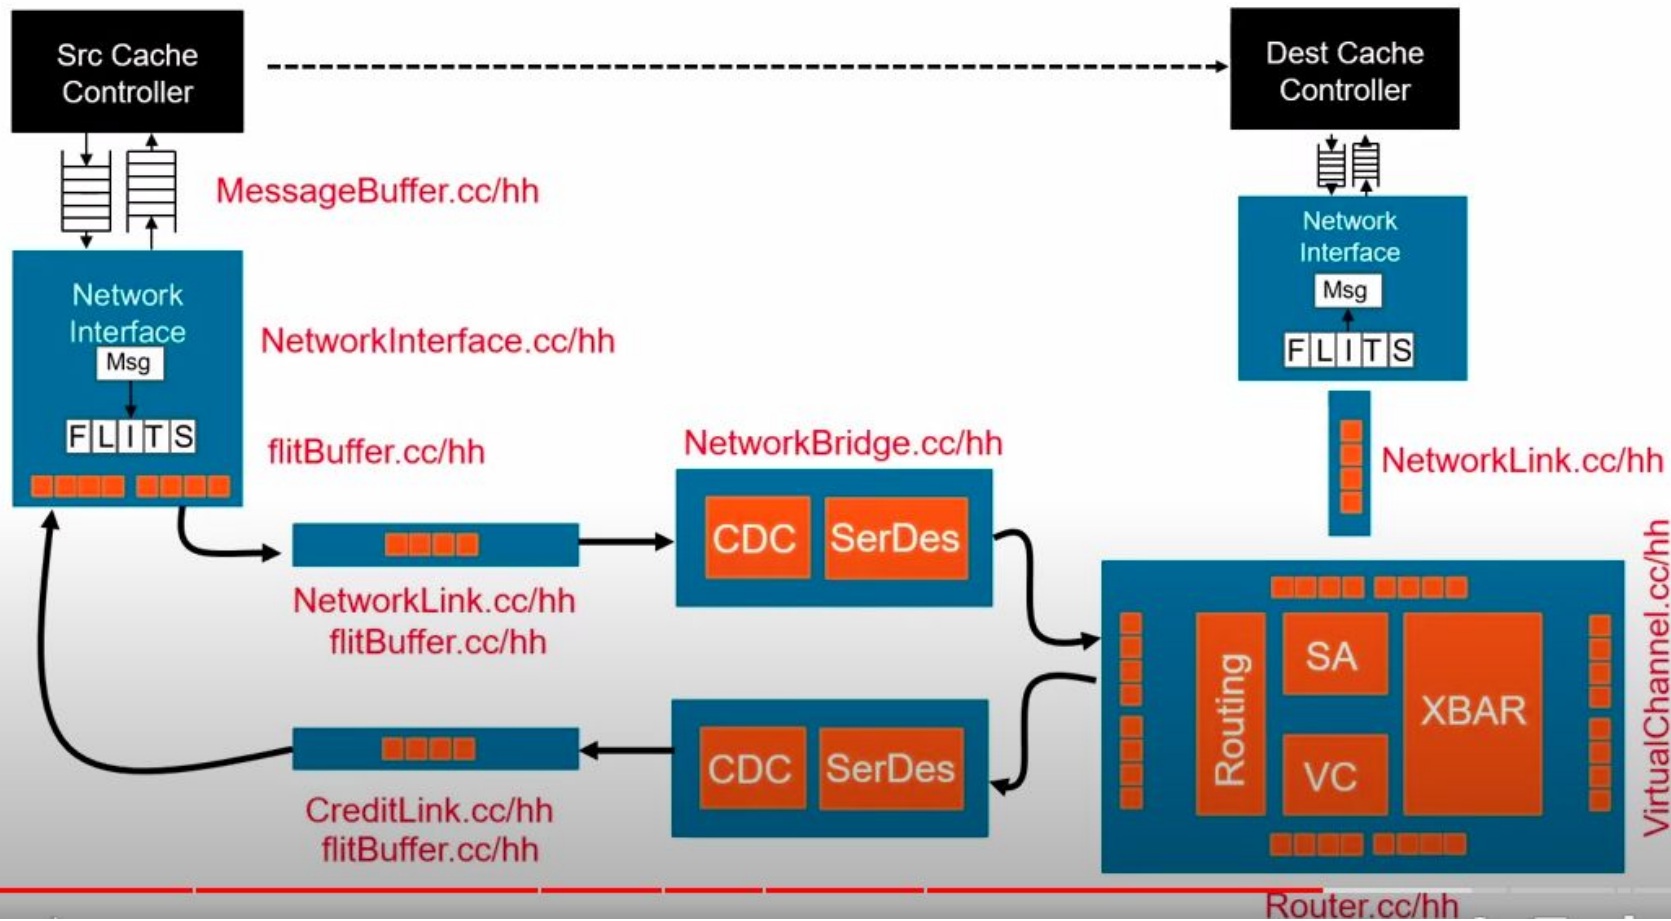


Serializer-Deserializer (SerDes)

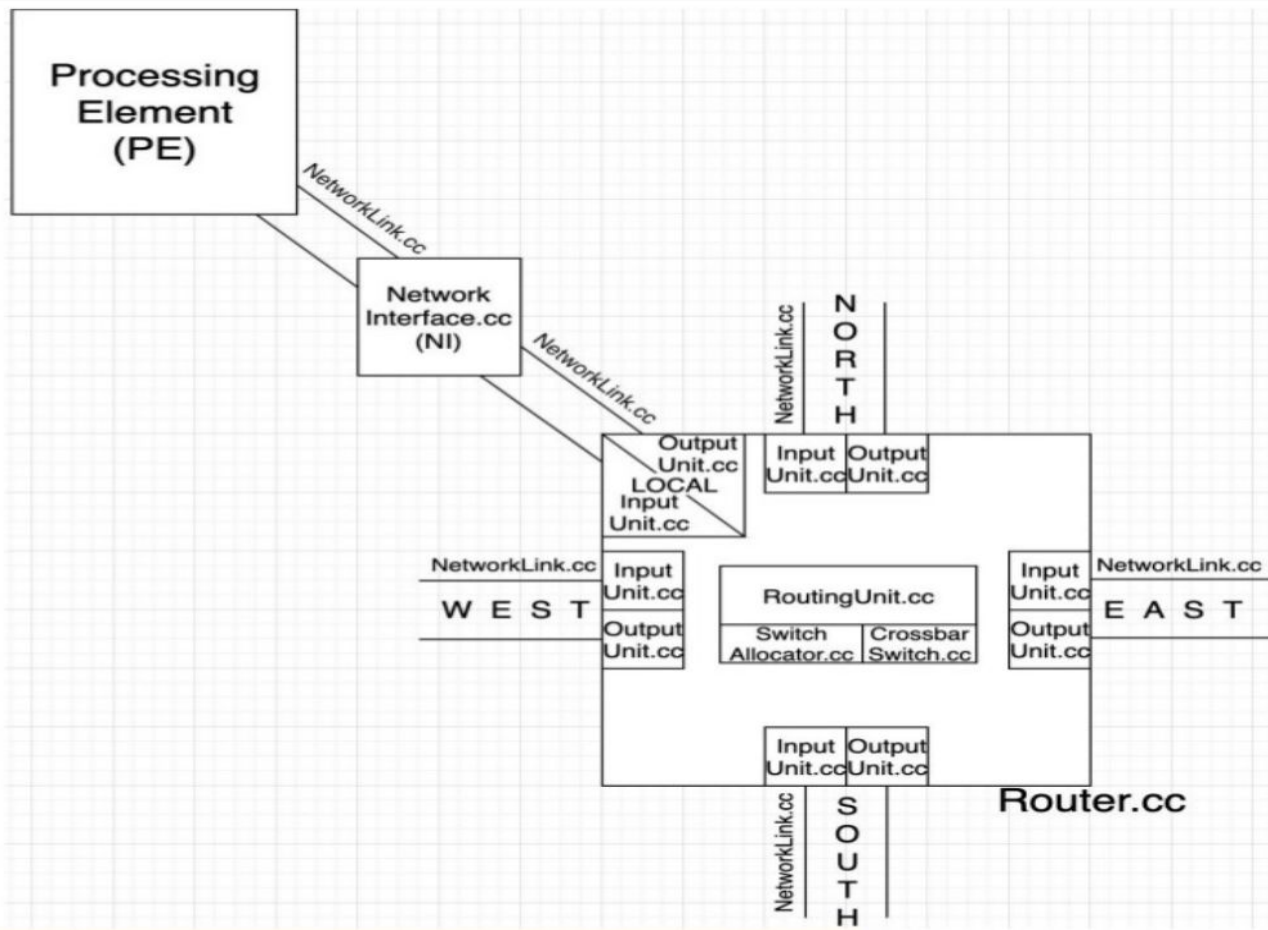
- Enables serialization or deserialization of flits
- Supports flow-control credit crossings
- Configurable latency
- Enabled by setting flags on the links
 - Internal Links: `src_serdes` and/or `dst_serdes`
 - External Links: `ext_serdes` and/or `int_serdes`




Life Cycle of a Message through HeteroGarnet



Garnet design: an overview



On-Chip Networks



Router

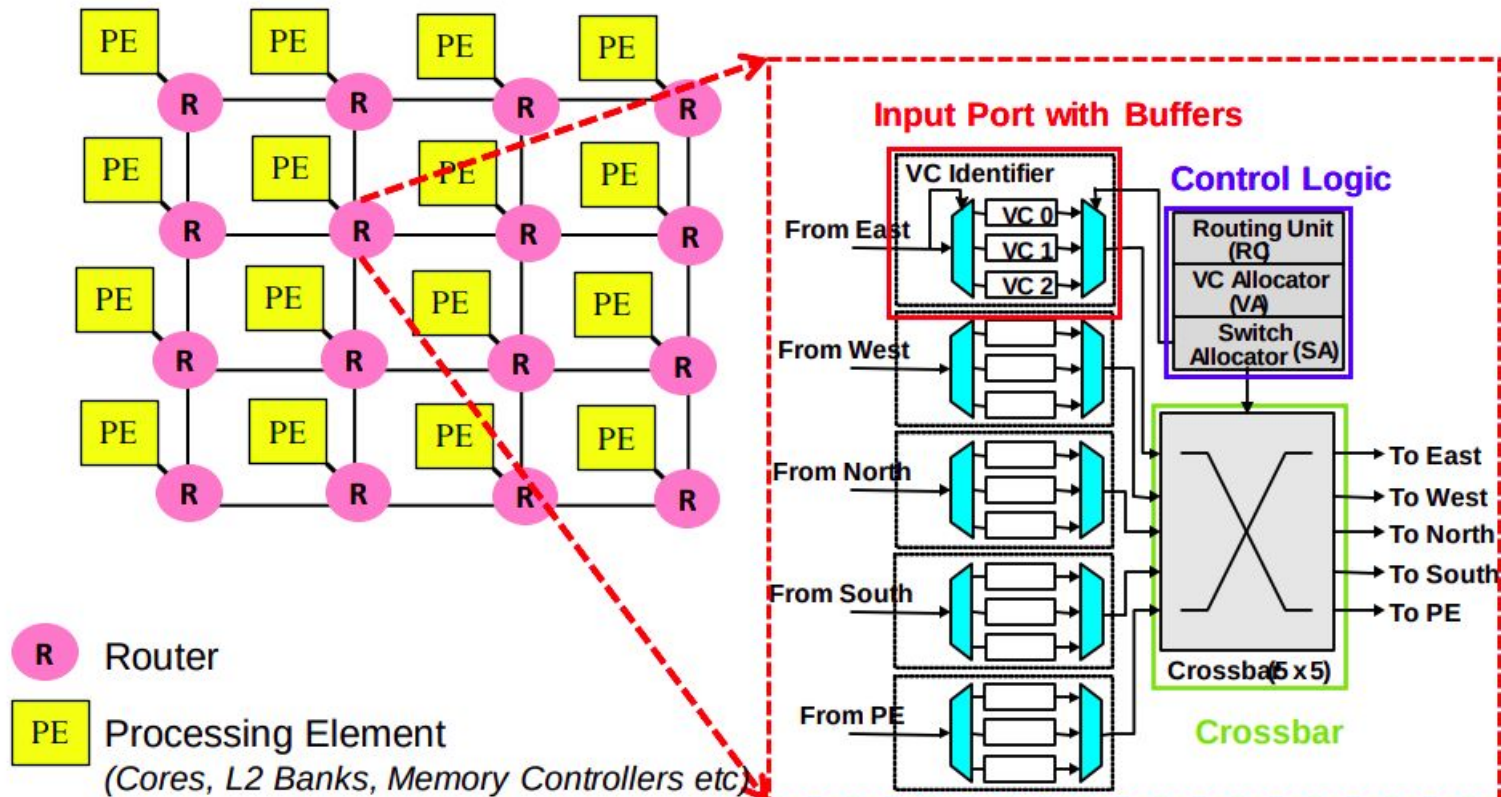


Processing Element

(Cores, L2 Banks, Memory Controllers, etc)

- Connect **cores, caches, memory controllers, etc**
 - Buses and crossbars are not scalable
- **Packet switched**
- **2D mesh:** Most commonly used topology
- Primarily serve **cache misses** and **memory requests**

On-chip Networks



Lifecycle of a Network Traversal

- `NetworkInterface.cc::wakeup()`
 - Every NI connected to one coherence protocol controller on one end, and one router on the other.
 - receives messages from coherence protocol buffer in appropriate vnet and converts them into network packets and sends them into the network.
 - garnet2.0 adds the ability to capture a network trace at this point [under development].
 - receives flits from the network, extracts the protocol message and sends it to the coherence protocol buffer in appropriate vnet.
 - manages flow-control (i.e., credits) with its attached router.
 - The consuming flit/credit output link of the NI is put in the global event queue with a timestamp set to next cycle. The eventqueue calls the wakeup function in the consumer.
- `NetworkLink.cc::wakeup()`
 - receives flits from NI/router and sends it to NI/router after `m_latency` cycles delay
 - Default latency value for every link can be set from command line (see `configs/network/Network.py`)
 - Per link latency can be overwritten in the topology file
 - The consumer of the link (NI/router) is put in the global event queue with a timestamp set after `m_latency` cycles. The eventqueue calls the wakeup function in the consumer.
- `Router.cc::wakeup()`
 - Loop through all `InputUnits` and call their `wakeup()`
 - Loop through all `OutputUnits` and call their `wakeup()`
 - Call `SwitchAllocator's` `wakeup()`
 - Call `CrossbarSwitch's` `wakeup()`
 - The router's wakeup function is called whenever any of its modules (`InputUnit`, `OutputUnit`, `SwitchAllocator`, `CrossbarSwitch`) have a ready flit/credit to act upon this cycle.

- InputUnit.cc::wakeup()
 - Read input flit from upstream router if it is ready for this cycle
 - For HEAD/HEAD_TAIL flits, perform route computation, and update route in the VC.
 - Buffer the flit for (m_latency - 1) cycles and mark it valid for SwitchAllocation starting that cycle.
 - Default latency for every router can be set from command line (see configs/network/Network.py)
 - Per router latency (i.e., num pipeline stages) can be set in the topology file.
- OutputUnit.cc::wakeup()
 - Read input credit from downstream router if it is ready for this cycle
 - Increment the credit in the appropriate output VC state.
 - Mark output VC as free if the credit carries is_free_signal as true
- SwitchAllocator.cc::wakeup()
 - Note: SwitchAllocator performs VC arbitration and selection within it.
 - SA-I (or SA-i): Loop through all input VCs at every input port, and select one in a round robin manner.
 - For HEAD/HEAD_TAIL flits only select an input VC whose output port has at least one free output VC.
 - For BODY/TAIL flits, only select an input VC that has credits in its output VC.
 - Place a request for the output port from this VC.
 - SA-II (or SA-o): Loop through all output ports, and select one input VC (that placed a request during SA-I) as the winner for this output port in a round robin manner.
 - For HEAD/HEAD_TAIL flits, perform outvc allocation (i.e., select a free VC from the output port.
 - For BODY/TAIL flits, decrement a credit in the output vc.
 - Read the flit out from the input VC, and send it to the CrossbarSwitch
 - Send an increment_credit signal to the upstream router for this input VC.
 - for HEAD_TAIL/TAIL flits, mark is_free_signal as true in the credit.
 - The input unit sends the credit out on the credit link to the upstream router.
 - Reschedule the Router to wakeup next cycle for any flits ready for SA next cycle.

- `CrossbarSwitch.cc::wakeup()`
 - Loop through all input ports, and send the winning flit out of its output port onto the output link.
 - The consuming flit output link of the router is put in the global event queue with a timestamp set to next cycle. The eventqueue calls the wakeup function in the consumer.
- `NetworkLink.cc::wakeup()`
 - receives flits from NI/router and sends it to NI/router after `m_latency` cycles delay
 - Default latency value for every link can be set from command line (see `configs/network/Network.py`)
 - Per link latency can be overwritten in the topology file
 - The consumer of the link (NI/router) is put in the global event queue with a timestamp set after `m_latency` cycles. The eventqueue calls the wakeup function in the consumer.

Thank you